Enhancing performance and reliability of Network File System

Aswin Babu Karuvally
Department of Computer Applications
College of Engineering Trivandrum, KTU
Trivandrum, Kerala
aswinbabuk@gmail.com

Basith Hameem
Department of Computer Applications
College of Engineering Trivandrum, KTU
Trivandrum, Kerala
basithhameem@cet.ac.in

Ann Jerin Sundar
Department of Computer Applications
College of Engineering Trivandrum, KTU
Trivandrum, Kerala
annjerinajs@gmail.com

Prof. John Prakash Joseph
Asst. Professor, Department of Computer Applications
College of Engineering Trivandrum, KTU
Trivandrum, Kerala
john@cet.ac.in

Abstract—**Network File System is a widely used distributed file system that allows the user to access and manipulate storage on remote computers, as if they were a part of the local machine. Network File System is notoriously slow in its default configuration and the incorporation of more than a dozen clients in the NFS environment merely accentuates the delay. When configured to deliver faster speeds by turning on asynchronous mode, the system suffers from higher risk of data corruption and loss.**

This study proposes variegated modifications to the Network File System, enabling it to provide elevated system performance, while containing the risk of data loss and corruption. Further, the proposed system behaves better in congested networks by consuming less bandwidth, ensuring decent speeds, even during periods of heavy network traffic.

Index Terms—**UNIX, NFS, Performance, Data loss, Data corruption**

I. INTRODUCTION

Network File System is a distributed File System protocol primarily used by the UNIX family of Operating Systems. It allows users to mount, access and manipulate disk partitions or directories on a remote computer, as if the said partition or directory was a part of the local machine. Network File System was developed as an open standard by SUN Microsystems in 1984.

NFS is widely used in Local Area Networks to conveniently share data, and provides users the ability to access their files across the network. Occasionally, a directory access protocol such as LDAP is combined with NFS, allowing the users to login to their user accounts from any computer on the network.

The crucial setback NFS depicts is the slow read and write speeds it offers with the default setup.The existent performance enhancing parameters in its default setup configuration files, either leave the performance rates unaffected or increases the probability of data corruption and loss. Thus the users are forced to run the system with the default, slow configuration.

The aforestated strategy diminishes the possibilities of interactive computing, and an Operating System requiring access to data on NFS share, often ends up freezing the computer in it's entirety, resulting environments with lost human productivity. Moreover, this bottlenecks the computer CPU, thereby exhausting valuable computing resources.

This paper proposes a number of changes to the Network File System protocol which increases the performance of Network File System while reducing the risk of data corruption and loss. The proposed system also ensures decent speeds in congested network as it consumes less bandwidth than the original NFS implementation and protects the ability of the NFS server to provide access to files during period of peak network activity.

II. BACKGROUND WORK

Past projects have proposed new Network File System protocols for transmitting data over Wide Area Networks. Some other papers have dealt with improving the performance of NFS over wireless links.

A.Muthitacharoen, B.Chen and D.Mazieres has proposed a Low Bandwidth Network File System, which can be used over Wide Area Networks such as the Internet. Though it minimizes the bandwidth usage of the protocol, it is meant for 90s era internet. To use it in current scenario would require extensive modification to the protocol and it is not compatible with SUN's NFS.

R.Dube, C.D Rais and S.K Tripathi has proposed a number of changes to the network stack to increase the performance of NFS over wireless links. The paper does not suggest changes to the NFS protocol, instead it tries to optimize the wireless network stack.

In our literature review, we are yet to come across a paper, which directly deals with enhancing the performance and

reducing data corruption rates of the widely adopted Network File System originally implemented by SUN Microsystems.

III. RESEARCH METHODOLOGY

All the benchmarks for this research paper has been done using NFS v4.2. The NFS environment comprising of client and server were emulated using virtual machines, with the Oracle VirtualBox hypervisor. Virtual Machines (VMs) provide a convinient method of simulating target hardware and networking infrastructure. The VirtualBox hypervisor was run on RedHat Enterprise Linux 7.2 with the hardware being an HP-Z640 workstation to ensure the VMs were not bottlenecked during benchmarks. The VMs were equipped with AMD PCnet-FAST III 100 Mbps network adapter, which was bridged to the network so that each VM was represented as a real machine on the network. Network Address Translation (NAT) was not used as NAT consumes CPU resources, which could in turn affect benchmarks. The NFS server and client VMs were set up to run Debian 9.3 (Stable) as their OS. Debian is a heavily tested distribution that ships with very stable packages. Debian was chosen so as to minimize the chance of bugs affecting the benhmarks.

The benchmarking tools used include Bonnie ++, Phoronix Test Suite/iozone, Phoronix Test Suite/dbench, Stand-alone Dbench and the "dd" utility. dd was used to perform benchmarks were client side cache effects needed to be avoided. Dbench is a powerful benchmarking tool designed to test the performance of NFS and SMB systems. It is capable of emulating upto 512 clients. All the benchmarks where done with Dbench emulating 6 clients. Benchmarking was often done through proper combination of all these tools. When eliminating client side cache effect was required, dd was run with "conv=fdatasync bs=1K" options to flush data from cache after write of each KB.

IV. WORKING AND TYPICAL PERFORMANCE OF NFS

Network File System is a distributed file system originally developed by SUN Microsystems in 1984. A distributed file system differ from normal file systems in the sense that they operate over the network and allow sharing of storage resources. NFS was developed as an open standard. It was initially implemented for UNIX but is now compatible with a wide array of Operating Systems. NFS has gone through four major revisions, with the first publicly available version being v2 (RFC 1094). All the experiments for this research have been conducted using NFS v4.2

NFS uses the Client-Server architecture. The NFS Server makes available a disk partition or directory on the network which can then be mounted just like a local storage device by the clients. For applications running on the client machine, NFS is just another file system and requires zero modifications inorder to work. This is made possible by an abstraction layer called Virtual File System or VFS. VFS defines what operations can be done on the filesystem regardless of the file system type. When an application deals with the UNIX file system, it is actually dealing with VFS. VFS receives data

regarding the target file from the application, then hands it over to the actual file system, which in this case is NFS.

Once NFS receives the data, it is transfered to the server with the help of Remote Procedure Call (RPC) in External Data Representation (XDR) format. RPC allows the NFS client to execute instructions on the server. XDR is a data representation standard that provides a uniform data format which can be understood by a variety of computers. This is one of the factors that provide NFS with cross platform compatibilty. The working of NFS protocol has been shown in figure.
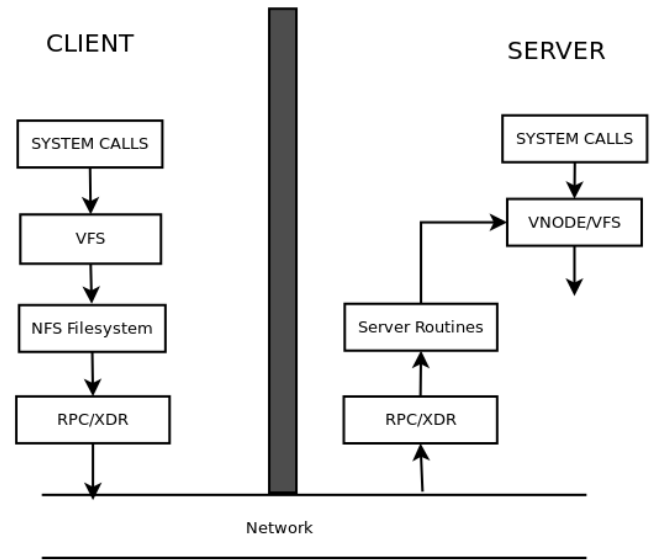

Fig. 1. VFS-NFS handover .

NFS by default runs in what is called server side synchronous mode. In this mode, when the NFS client receives a write operation, it connects to the server, requsts a write and transfers the data. Once the transfer is ;complete, the server syncs the data to its disks. After completing the sync operation, server returns an acknowledgment message back to the client. The issue here is that, the client has to wait till it receives the acknowledgment.It cannot perform any additional write to the server till the acknowledgment arrives. This considerably slows down the system. Further, clients are often configured to work in client side synchronous mode. In this mode, the client is forced to write the data to the server as soon as it receives the request. This often causes the system to crawl.

The performance of an NFS system with client and server side synchronous mode turned on was benchmarked with client side options in /etc/fstab set as [rw, sync, hard, intr 0 0] and the server side options in /etc/exports set as [rw, no-root-squash, subtree-check]. The benchmarks were conducted using dbench utility part of Phoronix Test Suite. This resulted in performance of 0.94 MB/s. The server and client systems in this case were equipped with 100mbps network adapters. The speeds obtained are thus clearly sub-optimal. The peformance graph obtained from the benchmark has been shown in figure.

V. PARAMETERS AFFECTING PERFORMANCE

NFS operates on the client-server architecture. There are tweaks that can be applied to NFS both at the client and server end. At the client end, the tweaks are made to /etc/fstab file and at the server end, the configuration is done to the /etc/exports file. Some of the commonly used options are listed below.

TABLE I
NFS OPTIONS-CLIENT SIDE

SlNo.	Option	Description
1	rw	Read/Write
2	syn	Sync file system with the server
3	hard	NFS requests are retried indefinitely
4	intr	Provided for backward compatibility
5	nfsvers	Specifies the nfs versions
6	rsize	Maximum number of bytes when reading data
7	wsize	Maximum number of bytes when writing data
8	udp	Specifies the connection to UDP
9	async	Asynchronous write

TABLE II
NFS OPTIONS-SERVER SIDE

SlNo.	Option	Description
1	rw	Read/Write
2	no-root-squash	Turn off root squashing
3	subtree-check	Specified directory/its subrectory for access
4	async	Synchronous write
5	sync	Asynchronous write

Various compbinations of these options were benchmarked. It was found that most of the options provided no increment in the performance of NFS environment with performance in the ball park of 0.94 MB/s.

Figure 0.77 MB/s

One notabble exception was the client side UDP option which reduced the speeds to 0.77 MB/s. At the end of experiments, it was found that considerable increase in performance was provided only by client side and server side async.

VI. PERFORMANCE WITH SERVER SIDE ASYNC

In server side synchronous mode, the server waits till the data has been written to its disk before returning the acknowledgment message to the client. Server side asynchronous mode changes the behaviour of NFS server such that it returns the acknowledgment message as soon as the client completes the transfer of data. This has tremendous impact on the performance of the system.

An NFS system with server side asynchronous mode was benchmarked with the client side options in /etc/fstab set as [rw, sync, hard, intr 0 0] and server side options in /etc/exports set as [rw,no-root-squash,no-subtree-check, async]. The test was conducted using dbench tool - part of phoronix test suite, and resulted in 28.72 MB/s, a huge increase in performance boost, compared with the earlier server side synchronous mode, which returned 0.94 MB/s. The performance graph obtained has been shown in figure. Further benchmarks were also conducted with dd utility, with a block size of 1KB, file

size of 2GB and fdatasync option turned on. The block size of 1K ensures the computer writes only 1KB of data at a time and the fdatasync option flushes the cache after each block is written. The aim was to understand the performance of the system without the contribution of client side caches. The benchmark resulted in 7.1 MB/s which was stil significantly higher than results recorded with server side asynchronous mode.Performance Figure. The higher performance comes from the fact that the clients do not have to wait till the server syncs the data with its disks. Clients can transfer data to the server, then get on with other tasks such as writing additional files. This also means that more number of clients can access the server in unit time. Still, with the current protocol, it is not advisable to leave server side async turned on due to the possibility of data loss and corruption.

A. Reliability concerns with Server side ASYNC

One side effect of enabling server side asynchronous mode is that, more clients can access the server in unit time. This in turn can cause a write queue to form on the client side. That is, writing a file to the server gives no guarantee that it has been written to permanent storage. In case the server crashes immediately after some data has been transferred to it, be it software crash or hardware failure, it can cause data corruption or loss. The more worrying fact is that, it is not just a single computer which looses data. Data loss can occour to most of the clients which has written to the server shortly before the crash.

If a server crash occours, a client has no means of protecting itself from the data loss. A client has no NFS cache that is permanent in nature. Even if the client has the lost file in its primary memory, there is a high chance of loosing it. This is because, if prior to crash the server was serving a critical configuration file, the application dependant on the file can crash. If the application is part of the Operating System, it can bring down the whole system. The latter is often the case with environments where home directory is served by an NFS share.

B. Fixing server side ASYNC behaviour

Once VFS handovers the write request to NFS client, it transfers the received data to the server rather than writing it to the local storage. In case of data loss, the file cannot be recovered, as the only copy of the file was in server's memory. The solution is to create a buffer in the client's local storage such that, a copy of all the data written to server will be kept with the client.

The buffer is a predefined storage area in the client's secondary memory. It acts like a ring buffer with a flexible memory size. The oldest files are deleted once the buffer reaches a predefined size. NFS client will maintain a plaintext file in the buffer containing names of each file in the buffer, its path and hash generated from each file. During an NFS write operation, the client stores a copy of the file in the buffer area and updates the metadata file with the information regarding the new file. The hash is calculated whenever a file is written to

the buffer. To minimize the CPU overhead for hash calculation, one should implement a lightweight hashing algorithm such as QUARK or PHOTON. Figure shows the working and structure of the buffer.

Figure of the working and structure of buffer + metadata file

Optional (Sample implementation of buffer deletion strategy)

In case of a server crash, the server creates a list of corrupted or lost files. During the first boot after the crash, the server requests the metadata file from each client that connects to it. Once the server receives the metadata, it calculates the hash for the local copy of each file that is listed in the metadata. If a file is missing or if the hashes do not match, they are added to retransmit-list, a list of files to be retransmitted from the client. Once the metadata file from a client is fully scanned, the retransmit list is sent to the client. The client in turn transmits a new copy of each file in the retransmit list. Figure shows the error recovery process.

C. Advantages of the proposed solution

By default the async mode in client side and server side is turned off. One of the main reason behind is the realiability and data loss. Concerning an organization data is very important. Loss of data is something that is difficult to accept. The above proposed solution cleanly avoid the data loss by implementing a circular buffer system. The advantages are:-

- Maintain high data speed.
- Eliminate data loss and corruption.
- The proposed system modifies only the NFS behaviour and leaves the underlying operating system untouched.
- End effect -¿ Leave server side async turned on.

VII. CLIENT SIDE ASYNCHRONOUS MODE

Client side asynchronous mode is another method, that can be used to improve the performance of Network File System. Client side asynchronous mode uses the client's RAM as a giant cache. Once enabled, writes to the NFS server are instead written to the client's RAM. The file gets written to server only if one of the following conditions are met:

- Memory fills up
- An application flushes cache using sync, fsync or msync
- A file is closed with close()
- A file is locked/unlocked using fcntl
- Client shutdown

That is, when using client side asynchronous mode, writes to the server are drastically reduced. This benefits both the client doing the write and the environment as a whole. For the client, it reduces the number of instances it has to transfer data to the server. For the server, it reduces the number of writes it has to deal with in unit time. This also means that the rest of the clients have to wait less to read/write to the server.

Client side asynchronous mode has mixed results from the benchmarks with pts/dbench measuring a marginally low 24 MB/s down from 28 MB/s with only server side asynchronous mode turned on , while dd measured 68 MB/s, an increase of 10 MB/s compared with only server side asynchronous mode turned on. Regardless of the benchmark, client side asynchronous mode certainly improves the performance of the network.

A. Enhancing client side ASYNC

The behaviour of client side asynchronous mode can be further enhanced to optimize the performance of an NFS environment. This can be done by modifying the server such that it keeps two variables, max-count and current-count. max-count is the maximum number of clients which can write to the NFS server at a time, while current-count is the number of clients which are currently writing to the server. max-count is a user configurable value and can be optimized as per the size of the network.

The client side asynchronous behaviour is modified such that, the clients try to write to the server as soon as a threshold is reached. The threshold is user configurable and can be optimized as per the configuration of the system. The server initially sets the value of current-count to 0. The server increments the value by one whenever a new client connects. When the client tries to connect to the server, the server checks if the current-count is less than max-count. The server allows the client to connect only if max-count - current-count ¿= 1. In case the server has reached its client limit, it denies the connection. The client waits for a random time rand-wait-time before trying to connect again. rand-wait-time is also user configurable and needs to be optimized per the size of the network. Figure shows the propsed behaviour of client side asynchronous mode

B. Advantages of proposed solution

The proposed solution provides a number of considerable advantages over the current system. First and foremost, it guards the data against data corruption. With the current system, a particular file can remain in the memory for too long. The proposed system forces the NFS client to flush its caches when a limit is reached. That is, a file won't remain in the cache for long, making sure it gets written to the server and to the proposed buffer, reducing the risk of data corruption.

The current client side asynchronous mode tries to fill up the RAM before trying to transfer the contents to the buffer. This can negatively affect the performance of both the applications running on the client and the network. If a particularly heavy application needs to be loaded, the client is forced to flush the caches, thus initiaing the NFS write. The client will need to wait for the write to complete before the application can be properly loaded. This introduces unnecessary latency to the system. The proposed system also makes sure that server's capacity is not wasted. The system makes sure a healthy number of clients gets to write to the server throughout its running time.

Regardless of the number of active clients, the number of clients doing write to the server will remain a constant. This can improve the overall performance of the environment by

ensuring that the server can provide decent read/write speeds regardless of the state of the environment. Thus the proposed system can provide a speedy and mature NFS behaviour.

C. Figures and Tables

a) Positioning Figures and Tables: Place figures and tables at the top and bottom of columns. Avoid placing them in the middle of columns. Large figures and tables may span across both columns. Figure captions should be below the figures; table heads should appear above the tables. Insert figures and tables after they are cited in the text. Use the abbreviation "Fig. 2", even at the beginning of a sentence.

TABLE III
NFS OPTIONS-CLIENT SIDE

| SlNo.	Option	Description
1	rw	Read/Write
2	syn	Sync file system with the server
3	hard	NFS requests are retried indefinitely
4	intr	Provided for backward compatibility
5	nfsvers	Specifies the nfs versions
6	rsize	Maximum number of bytes when reading data
7	wsize	Maximum number of bytes when writing data
8	udp	Specifies the connection to UDP
9	async	Asynchronous write

Fig. 2. Example of a figure caption.

Figure Labels: Use 8 point Times New Roman for Figure labels. Use words rather than symbols or abbreviations when writing Figure axis labels to avoid confusing the reader. As an example, write the quantity "Magnetization", or "Magnetization, M", not just "M". If including units in the label, present them within parentheses. Do not label axes only with units. In the example, write "Magnetization (A/m)" or "Magnetization {A[m(1)]}", not just "A/m". Do not label axes with a ratio of quantities and units. For example, write "Temperature (K)", not "Temperature/K".

REFERENCES

Please number citations consecutively within brackets [1]. The sentence punctuation follows the bracket [2]. Refer simply to the reference number, as in [3]—do not use "Ref. [3]" or "reference [3]" except at the beginning of a sentence: "Reference [3] was the first . . ."

Number footnotes separately in superscripts. Place the actual footnote at the bottom of the column in which it was cited. Do not put footnotes in the abstract or reference list. Use letters for table footnotes.

Unless there are six authors or more give all authors' names; do not use "et al.". Papers that have not been published, even if they have been submitted for publication, should be cited as "unpublished" [4]. Papers that have been accepted for publication should be cited as "in press" [5]. Capitalize only the first word in a paper title, except for proper nouns and element symbols.

For papers published in translation journals, please give the English citation first, followed by the original foreign-language citation [6].

REFERENCES

[1] G. Eason, B. Noble, and I. N. Sneddon, "On certain integrals of Lipschitz-Hankel type involving products of Bessel functions," Phil. Trans. Roy. Soc. London, vol. A247, pp. 529–551, April 1955.

[2] J. Clerk Maxwell, A Treatise on Electricity and Magnetism, 3rd ed., vol. 2. Oxford: Clarendon, 1892, pp.68–73.

[3] I. S. Jacobs and C. P. Bean, "Fine particles, thin films and exchange anisotropy," in Magnetism, vol. III, G. T. Rado and H. Suhl, Eds. New York: Academic, 1963, pp. 271–350.

[4] K. Elissa, "Title of paper if known," unpublished.

[5] R. Nicole, "Title of paper with only first word capitalized," J. Name Stand. Abbrev., in press.

[6] Y. Yorozu, M. Hirano, K. Oka, and Y. Tagawa, "Electron spectroscopy studies on magneto-optical media and plastic substrate interface," IEEE Transl. J. Magn. Japan, vol. 2, pp. 740–741, August 1987 [Digests 9th Annual Conf. Magnetics Japan, p. 301, 1982].

[7] M. Young, The Technical Writer's Handbook. Mill Valley, CA: University Science, 1989.